

Mock Demonstration Evaluation – AI-102 (Abhinav)

26 February 2026 07:13 PM

Evaluation Report

Date: 26 February 2026

Time: 19:20 (IST)

Subject: Mock Demonstration Analysis

Purpose: Trainer Readiness Assessment Based on Shared Transcript

Communication and Delivery Quality

1) Clarity of speech and articulation

- **What went wrong:** Key technical words (e.g., *efficiency*, *accuracy*) were under-articulated; initial vowel sounds were dropped and syllables blurred.
- **Why it matters:** Unclear diction increases cognitive load for learners, especially in diverse, multilingual classrooms, and risks mishearing critical terms.
- **How to correct:**
 - Practise target words using slow-then-fast drills and record-review.
 - Enunciate vowel onsets and hard consonants; keep the jaw relaxed and pace deliberate on technical terms.
- **Trainer-grade rewrite:** "...*efficiency* (ef-fi-shən-see) and *accuracy* (ak-yur-ə-see) are our two quality signals here... **slow down for the first syllable** ...notice how they show up in the model leaderboards we'll use in a minute."

2) Pacing and filler language

- **What went wrong:** Frequent "give me a second... okay... yes..." during navigation; "**let me think over it**" when asked a concept question.
- **Why it matters:** Filler and hedging signal uncertainty and reduce perceived control of the session.
- **How to correct:** Replace hedges with confident defer-and-park: "**Great question—coming up in two slides; I'll circle back after the demo.**" Keep silent micro-pauses instead of verbalised fillers.
- **Trainer-grade rewrite:** "Great question on LLMs. **We'll address that in the next slide;** for now, hold the thought and watch how this concept shows up in the model comparison."

3) Transitions and verbal presence

- **What went wrong:** Moves between slides and into the lab lacked preview/closure; navigation comments became stream-of-consciousness.
- **Why it matters:** Learners need signposts to maintain mental models.
- **How to correct:** Use a 3-beat transition—Close → Preview → Success criteria.
- **Trainer-grade rewrite:** "We've defined AI and positioned generative AI inside it. **Next**, we'll contrast **LLMs vs SLMs**. **Success today** is being able to pick the right model for a scenario... **this transition needs more structure** ...let's proceed."

4) On-screen writing legibility

- **What went wrong:** Pen-written terms ('input', 'prediction') were unclear; a few misspellings.

- **Why it matters:** Screenshots live on after class; unclear writing propagates confusion.
- **How to correct:** Pre-build typed shapes for formulas; if inking, zoom-in and print letters slowly.
- **Trainer-grade rewrite:** “I’ll **type** the mapping: $y = f(x)$ — y is output, x is input, and f is the learned function... **here, slow down** ...we’ll use this mapping again in the lab.”

Conceptual Accuracy and Technical Validity

1) Function mapping language

- **What went wrong:** Saying the function “gets **multiplied** into” the input is mathematically inaccurate for a general learned mapping.
- **Why it matters:** Implies linearity and misguides learners about model behaviour.
- **How to correct:** “ $y = f(x)$ reads ‘ f maps input x to output y ’; f is learned from data.”
- **Trainer-grade rewrite:** “Think of f as the **learned rule**: given x (customer features), the model returns y (fraud likelihood). No multiplication implied.”

2) LLM vs SLM framing

- **What went right:** Correct directionality—LLMs (e.g., GPT-4 class) vs SLMs (e.g., Microsoft Phi family).
- **What to strengthen:** Anchor with Microsoft’s positioning: SLMs optimise for speed/footprint; LLMs for breadth/quality—show both in Foundry’s catalog and leaderboards. [\[azure.microsoft.com\]](https://azure.microsoft.com)
- **Trainer-grade rewrite:** “**LLMs** (e.g., GPT-4o) maximise quality and multimodality; **SLMs** (e.g., **Phi-4**) trade some breadth for cost and latency—ideal for edge or constrained workloads.” [\[azure.microsoft.com\]](https://azure.microsoft.com), [\[ai.azure.com\]](https://ai.azure.com)

3) Model catalog metrics (quality, safety, cost, throughput)

- **What went right:** You named the metrics during comparison—good.
- **What to strengthen:** Explicitly reference **Model Leaderboards** and **trade-off charts** so learners can self-serve later. [\[learn.microsoft.com\]](https://learn.microsoft.com), [\[learn.microsoft.com\]](https://learn.microsoft.com)
- **Trainer-grade rewrite:** “In **Microsoft Foundry** → **Model Catalog** → **Compare models**, the **leaderboards** summarise **quality, safety, cost, throughput**; the **trade-off chart** lets us pick the best fit for our scenario.” [\[learn.microsoft.com\]](https://learn.microsoft.com)

4) Platform naming and currency

- **What went wrong:** Mixed usage of “Azure AI Foundry” and “Microsoft AI Foundry.”
- **Why it matters:** Current branding is **Microsoft Foundry**; consistency builds authority and prevents confusion. [\[schneider.im\]](https://schneider.im)
- **Trainer-grade rewrite:** “We’ll use **Microsoft Foundry** (formerly Azure AI Foundry) to explore the model catalog and deploy quickly.” [\[schneider.im\]](https://schneider.im)

5) RAG context

- **What went wrong:** RAG was mentioned but not grounded in an architecture.
- **Why it matters:** AI-102 learners expect to see retrieval (Azure AI Search), grounding, and generation in one flow; it reduces hallucination and aligns with Microsoft Learn guidance. [\[learn.microsoft.com\]](https://learn.microsoft.com)
- **How to correct:** Add a one-slide bird’s-eye diagram: **Query** → **Orchestrator** → **Azure AI Search (retrieve)** → **Prompt with citations** → **LLM** → **Answer + sources**; note classic vs agentic retrieval choices. [\[learn.microsoft.com\]](https://learn.microsoft.com)

- **Trainer-grade rewrite:** “At a glance—**RAG** grounds responses on your content: **retrieve** with Azure AI Search, **augment** the prompt, and **generate** with an LLM; choose **classic** or **agentic retrieval** based on complexity.” [\[learn.microsoft.com\]](https://learn.microsoft.com)

Structural Flow and Teaching Logic

1) Topic introductions & context build-up

- **Observation:** Good on AI vs GenAI context; missing explicit learning outcomes.
- **Fix:** Start with 3 *outcomes* (“select a model, explain LLM vs SLM, compare GPT-4o vs Phi-4 for a scenario”). Tie each section back to these outcomes.
- **Trainer-grade rewrite:** “Today’s outcomes: (1) Navigate Foundry’s catalog and **compare models**; (2) Explain **LLM vs SLM** trade-offs; (3) **Choose** GPT-4o or Phi-4 for a given workload.” [\[learn.microsoft.com\]](https://learn.microsoft.com)

2) Simple → complex progression

- **Observation:** Reasonable concept order; the lab jump happened mid-concept with no synthesis.
- **Fix:** Add a 30-second recap and a clear lab objective.
- **Trainer-grade rewrite:** “Quick recap: GenAI creates; LLMs vs SLMs define quality vs cost. **Lab objective:** open **Model Catalog**, compare **GPT-4o** vs **Phi-4-reasoning**, and pick one for a help-desk bot.” [\[microsoft.github.io\]](https://microsoft.github.io)

3) Summarisation & conceptual closure

- **Observation:** Section endings were soft.
- **Fix:** Close each block with 3 key points + 1 check question to consolidate learning.
- **Trainer-grade rewrite:** “**Summary:** LLM vs SLM, trade-offs, and when to use each. **Checkpoint:** Which metric matters most for a high-traffic chat? Throughput—why?” [\[learn.microsoft.com\]](https://learn.microsoft.com)

Classroom Intelligence and Learner Orientation

1) Anticipation of confusion

- **Observation:** Referring to “this figure” without a visible pointer left learners searching.
- **Fix:** Always pair verbal reference with a laser or callout; name the region (“left pane, green bubble”).
- **Trainer-grade rewrite:** “Focus on the **left pane—Model Leaderboards**, the **green tiles** show quality and safety at a glance...” [\[learn.microsoft.com\]](https://learn.microsoft.com)

2) Analogy and simplification

- **Opportunity:** Use a toolbox analogy for model selection (power tool vs precision screwdriver) to make LLM/SLM trade-offs tangible.
- **Trainer-grade rewrite:** “Think of **GPT-4o** as a versatile **power tool**—heavier but powerful; **Phi-4** as the **precision screwdriver**—lighter, faster, cheaper for focused tasks.” [\[azure.microsoft.com\]](https://azure.microsoft.com), [\[ai.azure.com\]](https://ai.azure.com)

3) Engagement orientation

- **Observation:** Polite “hope you understood” checks; consider structured micro-checks (polls, thumbs, or one-line chat answers) every 6–8 minutes.

Real-World and Industry Context

- **Observation:** Enterprise ties were light.
- **Fix:** Offer concrete vignettes mapping metrics to scenarios:
 - **Customer chat at scale:** prioritise **throughput & cost** → trial **Phi-4-reasoning** first; if quality gaps appear, escalate to **GPT-4o**. [\[azure.microsoft.com\]](https://azure.microsoft.com), [\[ai.azure.com\]](https://ai.azure.com)
 - **Vision-augmented triage:** need **multimodal quality** → **GPT-4o** with image inputs. [\[ai.azure.com\]](https://ai.azure.com)
 - **Edge/latency-sensitive tasks:** **SLM (Phi-4)** for local/low-latency deployments. [\[azure.microsoft.com\]](https://azure.microsoft.com)

Slide and Visual Navigation Accuracy

- **What went wrong:** No active laser when referencing figures; instruction pane hidden in the lab.
- **Why it matters:** Learners lose context and can't mirror steps.
- **How to correct:**
 - Keep **instructions visible** (minimised rather than closed) when demonstrating a hosted lab, or paste the relevant **Microsoft Learn** exercise link in chat for self-paced following. [\[microsoftl...github.io\]](https://microsoftl...github.io)
 - Use a pointer and verbal coordinates on every figure reference.
- **Trainer-grade rewrite:** *"I'll keep the **instruction pane** docked on the right while I execute steps on the left—you can follow the same **MS Learn** link in chat."* [\[microsoftl...github.io\]](https://microsoftl...github.io)

Behavioural Maturity and Professional Tone

- **Composure & patience:** Good steadiness during portal slowness; you negotiated alternatives calmly.
- **Openness & self-awareness:** Willingness to accept suggestions (e.g., showing instructions, adding architecture overview).
- **Refinement cue:** Replace apology loops with **structured fallback** language: *"If this page doesn't load in 30 seconds, we'll switch to the step-by-step deck and resume live after the break."*

Question-Handling Effectiveness

- **Observation:** Deferral of LLM question was reasonable; phrasing ("let me think over it") signalled uncertainty.
- **Fix:** **Micro-answer + pointer** or **defer-and-park** with time-box.
- **Trainer-grade rewrite:** *"Large Language Models—pretrained models with billions of parameters that generalise across tasks—we'll **unpack this next** and compare to smaller **Phi** models for cost/latency."* [\[azure.microsoft.com\]](https://azure.microsoft.com)

Error-Handling Discipline

- **Observation:** You attempted relaunch and then pivoted to narration—good.

- **Fix:** Maintain a **pre-built screenshot deck** and a 90-second silent screen capture for the core workflow (Model Catalog → Compare → Playground) to cover service latency days. Also maintain **consistent naming** (Microsoft Foundry) during plan-B pivots. [\[schneider.im\]](#)

Long-Term Training Potential

- **Assessment:** You're coachable, conceptually grounded at a high level, and improving in pace and tone. With targeted work on **terminology precision**, **visual signalling**, and **lab hygiene**, you can deliver AI-102 content with confidence.

Risk Assessment

- **Knowledge gaps:** Limited hands-on time visible in **Microsoft Foundry** flows; schedule two end-to-end dry runs (catalog → compare → deploy → evaluate). [\[learn.microsoft.com\]](#)
- **Structural weaknesses:** Soft transitions and light section closures.
- **Unclear explanations:** Ambiguous mapping language for $y = f(x)$.
- **Engagement issues:** Minimal structured checks for understanding.
- **Confidence instability:** Hedging phrases reduce authority.
- **Credibility risks:** Branding inconsistency (Azure AI Foundry vs Microsoft Foundry). [\[schneider.im\]](#)

Ideal Delivery (Concise 10–15 Line Rewrite)

"Welcome—today we'll achieve three outcomes: navigate **Microsoft Foundry**, explain **LLM vs SLM** trade-offs, and **choose** a model for a real scenario. **First**, AI vs GenAI: AI analyses; **GenAI** creates. $y = f(x)$ means ' f maps input x to output y '—no multiplication implied... **this point benefits from a practical scenario** ...customer data in, fraud likelihood out. **LLM vs SLM:** **GPT-4o** for high-quality multimodal tasks; **Phi-4** for low-latency, cost-efficient tasks or edge. **Now the architecture big picture:** user → orchestrator → **Azure AI Search (retrieve)** → prompt grounding → **LLM** → answer with citations. **In Foundry's Model Catalog**, we'll compare **GPT-4o** and **Phi-4-reasoning** on **quality, safety, cost, throughput** using the **leaderboards** and **trade-off chart**... **this transition needs more structure** ...we'll pick based on our help-desk use case. **Lab plan:** instructions on the right, actions on the left; if the portal lags, we'll use the step deck and resume live later. **Checkpoint:** which metric dominates a high-traffic chat? **Throughput** first, then cost. **Wrap-up:** you can now select a model confidently, explain the trade-offs, and justify your choice." [\[learn.microsoft.com\]](#), [\[learn.microsoft.com\]](#), [\[learn.microsoft.com\]](#), [\[schneider.im\]](#)

A. Feedback Summary

Strengths

- Improved pacing and a more natural, conversational tone.
- Correct high-level framing of **LLM vs SLM** and intent to use Foundry comparison. [\[azure.microsoft.com\]](#), [\[learn.microsoft.com\]](#)
- Composure under platform slowness; willingness to pivot.

Weaknesses

- **Articulation and on-screen spelling;** diagram references without visual anchoring.
- **Branding inconsistency** and **soft transitions** with limited section closure. [\[schneider.im\]](#)

- RAG left unframed; minimal structured checks for understanding. [\[learn.microsoft.com\]](https://learn.microsoft.com)

Behavioural Patterns

- Polite, receptive, and steady under pressure; occasional hedging that dilutes authority.

Conceptual Clarity

- Solid mental model; needs precision in mathematical phrasing and platform terminology.

B. Remarks

You've made visible progress on pace and tone. Focus next on **precision** (mathematical language and platform naming), **visual discipline** (pointer/callouts every time), and **lab hygiene** (instructions visible, links shared). Open with an **architecture overview** and close each block with **three takeaways + a check**; this alone will elevate classroom clarity and confidence.

[\[learn.microsoft.com\]](https://learn.microsoft.com), [\[learn.microsoft.com\]](https://learn.microsoft.com)

C. Final Verdict

You show **growing confidence**, **coachability**, and a **workable mental model** of GenAI delivery. To be fully classroom-ready, tighten **terminology**, **transitions**, and **visual signalling**, and rehearse an end-to-end **Microsoft Foundry** flow with the leaderboards/trade-off comparison and a brief RAG overview. With these refinements, learner impact will rise sharply and Q&A will feel effortless. At present, you are **close**, but require a short cycle of focused practice and a final validation run.

Rating: Satisfactory

Readiness Category: Needs Moderate Improvement